

2. Round-off and truncation errors


- significant figures
- numerical errors
- round-off errors
- Taylor series, truncation errors
- error propagation

Significant figures

significant figures indicate the reliability of a numerical value

- represent certain digits plus one estimated digit
- ensure confidence in computations

Example



speedometer (48.5 km/h, 3 significant figures with 2 certain digits)

odometer (87,324.45 km, 7 significant figures)

Rules for significant figures

- non-zero digits are always significant
- zeros between non-zero digits are significant
 - e.g., 1002 has 4 significant figures
- leading zeros are not significant
 - e.g., 0.001845 has 4 significant figures
- trailing zeros with a decimal point are significant
 - e.g., 45.300 has 5 significant figures
- trailing zeros can be significant or not
 - e.g., 45,300 may have 3, 4, 5 significant figures
 - use scientific notation for clarity
 - e.g., 4.5300×10^4 has 5 significant figures
- computers retain only a finite number of significant figures
 - e.g., $\pi = 3.141592653589793238462643\cdots$
 - π cannot be represented exactly in a computer
 - omission of the remaining significant figures is called *round-off error*


Accuracy and precision

Accuracy: how close a value is to the true value

Precision: how close repeated values are to each other

Example

- (a) inaccurate (biased), imprecise (uncertain)
- (b) accurate, imprecise
- (c) inaccurate, precise
- (d) accurate, precise



Outline

- significant figures
- **numerical errors**
- round-off errors
- Taylor series, truncation errors
- error propagation

Error sources

Errors in the problem to be solved

- mathematical model errors (model approximation)
- error in input data (measurement and previous approximate computation)

Truncation and discretization errors

- due to using approximate formula
 - replacing derivatives by finite differences
 - evaluating function by truncating a Taylor series (e.g., $e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots$)
- *convergence errors* in iterative methods, which converge to the exact solution in infinitely many iterations, but are cut off after a finite number of iterations

Round-off errors

- arise from finite precision representation of real numbers on computers
- truncation or discretization errors usually dominate round-off errors

Example

surface area of the Earth with radius r might be computed using the formula

$$A = 4\pi r^2$$

- earth is modeled as a sphere, which is an approximation of its true shape
- $r \approx 6370$ km, is based on empirical measurements and previous computations
- π is given by an infinite limiting process, which must be truncated at some point
- numerical values for the input data, as well as the results of the arithmetic operations performed on them, are rounded in a computer or calculator

Approximations and errors

given true (actual) value x and its approximation \hat{x}

- *true error:* $E_t = x - \hat{x}$ (*absolute error:* $|x - \hat{x}|$)
- *relative error:*

$$\frac{|\text{true error}|}{|\text{true value}|} = \frac{|x - \hat{x}|}{|x|} \quad (\text{assuming } x \neq 0)$$

- gives percentage of error compared to the actual value ($\varepsilon_t = \text{relative error} \times 100\%$)
- accounts of the order of magnitude of quantities

Example

x	\hat{x}	absolute error	relative error
1	0.99	0.01	0.01
1	1.01	0.01	0.01
100	99.99	0.01	0.0001
100	99	1	0.01

- when $|x| \approx 1$, little difference between absolute and relative error
- when $|x| \gg 1$, relative error more meaningful

Example: calculation of errors

Problem

- (a) measurement of a bridge 9999 cm with true value 10,000 cm
- (b) measurement of a rivet 9 cm with true value 10 cm

Error

- (a) bridge: $E_t = 10,000 - 9999 = 1 \text{ cm}$ and $\varepsilon_t = \frac{1}{10,000} \times 100\% = 0.01\%$
- (b) rivet: $E_t = 10 - 9 = 1 \text{ cm}$ and $\varepsilon_t = \frac{1}{10} \times 100\% = 10\%$

- same error (1 cm) but different relative impact
- conclusion: we did a good job of measuring the bridge, but not the rivet

Approximate error

- true value is typically unknown a priori
- use best available estimate

$$\varepsilon_a = \frac{\text{approximate error}}{\text{approximation}} \times 100\% = \frac{x_i - x_{i-1}}{x_i} \times 100\%$$

where x_i is current approximation and x_{i-1} is previous approximation

- typically we require $|\varepsilon_a| < \varepsilon_s$ for some small tolerance $\varepsilon_s > 0$
 - referred to as a *stopping criterion*
- we can be assured that the result is correct to *at least n* significant figures if

$$\varepsilon_s = (0.5 \times 10^{2-n})\%$$

Example: error estimates for iterative methods

Problem: estimate $e^{0.5} = 1.648721\cdots$ using *Maclaurin series*

$$e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \cdots$$

Goal: achieve 3 significant figures $\Rightarrow \varepsilon_s = 0.5 \times 10^{2-3}\% = 0.05\%$

- add terms until $|\varepsilon_a| < \varepsilon_s$
- example: for $x = 0.5$, first term (1), second term (1.5), etc.

terms	$e^{0.5} \approx$	$\varepsilon_t (\%)$	$\varepsilon_a (\%)$
1	1	39.3	
2	1.5	9.02	33.3
3	1.625	1.44	7.69
4	1.645833333	0.175	1.27
5	1.648437500	0.0172	0.158
6	1.648697917	0.00142	0.0158

- stops at 6 terms: $\varepsilon_a < 0.05\%$, accurate to 5 significant figures

Example: MATLAB iterative calculation

implement code to iteratively find $e^x \approx \sum_{i=0}^n \frac{x^i}{i!}$ for $x = 1$ until $|\varepsilon_a| < 10^{-6}$

```
>> format long
>> [val, ea, iter] = IterMeth(1,1e-6,100)
val =
2.718281826198493
ea =
9.216155641522974e-007
iter =
12

function [fx,ea,iter] = IterMeth(x,es,maxit)
if nargin < 2 || isempty(es),es = 0.0001;end
if nargin < 3 || isempty(maxit),maxit = 50;end
% initialization
iter = 1; sol = 1; ea = 100;
% iterative calculation
while (1)
solold = sol;
sol = sol + x^iter / factorial(iter);
iter = iter + 1;
if sol ~= 0
ea = abs((sol - solold)/sol)*100;
end
if ea<= es || iter>= maxit,break,end
end
fx = sol;
end
numerical errors
```

Trade-offs in numerical methods

Accuracy vs. efficiency

- more accurate methods increase computation time
- e.g., numerical integration with more intervals improves accs. but slows comp.

Stability

- some methods become unstable for certain problems
- e.g., stiff differential equations

Convergence: does the method approach the true solution as iterations increase?

Outline

- significant figures
- numerical errors
- **round-off errors**
- Taylor series, truncation errors
- error propagation

Round-off errors

- computers retain only a fixed number of significant figures
- irrational no. (π , e , $\sqrt{7}$) and man rationals cannot be represented exactly
- *round-off error*: discrepancy introduced by omitting significant figures

Example: $\sqrt{2} = 1.4142135623731\dots \approx 1.4142$ using 5 significant figures

- $(\sqrt{2})^5 = 5.6569$ and $(1.4142)^5 = 5.6566$ (small difference)
- $(\sqrt{2})^{50} = 33554432$ and $(1.4142)^{50} = 33538346.35$ (big difference ≈ 1600)

Computer number systems

Base-10 (decimal)

- uses digits 0–9 with place values 10^k
- example: $86,409 = 8 \times 10^4 + 6 \times 10^3 + 4 \times 10^2 + 0 \times 10^1 + 9 \times 10^0$
- example: $8642.9 = 8 \times 10^3 + 6 \times 10^2 + 4 \times 10^1 + 2 \times 10^0 + 9 \times 10^{-1}$


Base-2 (binary)

- uses digits 0, 1 with place values 2^k
- example: $(10101101)_2 = 1 \cdot 2^7 + 0 \cdot 2^6 + \dots + 1 \cdot 2^0 = 173$
- example: $(101.1)_2 = 1 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 + 1 \times 2^{-1} = 5.5$

Integer representation

Integer representation

- *signed magnitude*: first bit is sign (0: +ve, 1: -ve); remaining bits store magnitude
- example: -173 (16-bit signed magnitude):



Example: determine the base-10 range on a 16-bit machine

- 1 bit for sign; 15 bits for magnitude: max unsigned magnitude

$$(1 \times 2^{14}) + (1 \times 2^{13}) + \cdots + (1 \times 2^1) + (1 \times 2^0) = 2^{15} - 1 = 32,767$$


- zero is 0000000000000000, so it is redundant to use 1000000000000000
- store additional -ve number: so range -2^{15} to $2^{15} - 1 = -32,768$ to 32,767
- numbers outside this range cannot be represented (**overflow/underflow**)

Floating-point representation

fractional numbers are represented in computers using *floating-point* form

$$x = \pm(.d_1 d_2 \cdots d_n) \cdot b^e = \pm \left(\frac{d_1}{b^1} + \cdots + \frac{d_n}{b^n} \right) \cdot b^e = \pm m \cdot b^e$$

- b is the *base* (an integer larger than 1); n is *precision* (number of digits)
- e is *exponent* ($e_{\min} \leq e \leq e_{\max}$)
- $d_1 d_2 d_3 \cdots$ is *mantissa* or *significand*, d_i integer with $0 \leq d_i \leq b - 1$
- stored as:



Examples

- base-10: $0.15678 \times 10^3 = 156.78$
- base-2: $-(.1101) \cdot 2^2 = -(\frac{1}{2} + \frac{1}{4} + \frac{0}{8} + \frac{1}{16}) \cdot 2^2 = -3.25$

Normalization of mantissa

to maximize significant figures, mantissa is **normalized** to remove leading zeros

- example:

$$1/34 = 0.029411765\dots$$

so

store as $0.0294 \times 10^0 \Rightarrow$ normalize to 0.2941×10^{-1}

- normalization bounds mantissa:


$$\frac{1}{b} \leq m < 1$$

$0.1 \leq m \leq 1$ for base-10 and for $0.5 \leq m \leq 1$ base-2

- for binary, we have $d_1 = 1$, which need not be stored

Example: binary 7-bit floating-point set

7-bit floating point number stored as



- smallest positive normalized value (shown above):

$$m = 1 \times 2^{-1} = 0.5, \quad e = -(1 \times 2^1 + 1 \times 2^0) = -3 \Rightarrow +0.5 \times 2^{-3} = 0.0625$$

- next highest numbers are developed by increasing the mantissa, as in

$$0111101 = (1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3}) \times 2^{-3} = (0.078125)_{10}$$

$$0111110 = (1 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3}) \times 2^{-3} = (0.093750)_{10}$$

$$0111111 = (1 \times 2^{-1} + 1 \times 2^{-2} + 1 \times 2^{-3}) \times 2^{-3} = (0.109375)_{10}$$

in base-10 equivalents are spaced evenly with an interval of 0.015625

Example: binary 7-bit floating-point set

- to continue increasing, we decrease the exponent to 10, which gives a value of

$$e = -(1 \times 2^1 + 0 \times 2^0) = -2$$

- mantissa is decreased back to its smallest 100; so, next number is

$$0110100 = (1 \times 2^{-1} + 0 \times 2^{-2} + 0 \times 2^{-3}) \times 2^{-2} = (0.125000)_{10}$$

this still represents a gap of $0.125000 - 0.109375 = 0.015625$

- increasing the mantissa, the gap is lengthened to 0.03125:

$$0110101 = (1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3}) \times 2^{-2} = (0.156250)_{10}$$


$$0110110 = (1 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3}) \times 2^{-2} = (0.187500)_{10}$$

$$0110111 = (1 \times 2^{-1} + 1 \times 2^{-2} + 1 \times 2^{-3}) \times 2^{-2} = (0.218750)_{10}$$

- this pattern is repeated until a maximum number is reached:

$$0011111 = (1 \times 2^{-1} + 1 \times 2^{-2} + 1 \times 2^{-3}) \times 2^3 = (7)_{10}$$

Floating-point consequences and errors



- limited range (overflow and underflow)
- only a finite number of numbers can be represented within the range
- the interval between numbers, Δx , doubles as the numbers grow in magnitude

Chopping and rounding

Chopping (truncation)

- discard excess digits (bias toward lower endpoint)
- example (base-10, 7 sig figs): $\pi = 3.141592653\cdots$

chop: 3.141592

Rounding

- map to nearest representable number (reduced error, unbiased overall)
- example (base-10, 7 sig figs): $\pi = 3.141592653\cdots$

round: 3.141593

Relative error bounds and machine epsilon

for $\Delta x = \text{actual number} - \text{floating point representation}$, we have

- for chopping:

$$\frac{|\Delta x|}{|x|} \leq \epsilon$$

- for rounding:

$$\frac{|\Delta x|}{|x|} \leq \frac{\epsilon}{2}$$


ϵ is **machine epsilon**

$$\epsilon = b^{1-n}$$

where b is base and n is mantissa digits (precision)

Example: machine epsilon for the 7-bit set

- base $b = 2$, mantissa bits $n = 3$ and chopping, we have $\mathcal{E} = 2^{1-3} = 0.25$
- largest error occurs just below the upper bound of the 1st normalized interval



- for example, maximum error would be a value falling just below the upper bound of the interval between $(0.125000)_{10}$ and $(0.156250)_{10}$
- for this case, the error is less than

$$\frac{|\Delta x|}{|x|} < \frac{0.03125}{0.125000} = 0.25$$

IEEE standard for binary arithmetic

- two binary ($b = 2$) floating-point number systems
- used in almost all modern computers (e.g., MATLAB)

IEEE standard single precision (requires 32 bits)

$$n = 24, \quad e_{\min} = -125, \quad e_{\max} = 128$$

- 23 bits for mantissa ($d_1 = 1$ not stored)
- 1 sign bit and 8 bits for exponent
- about 7 significant base-10 digits precision with range 10^{-38} to 10^{39}

IEEE standard double precision (requires 64 bits)

$$n = 53, \quad e_{\min} = -1021, \quad e_{\max} = 1024$$

- 52 bits for mantissa ($d_1 = 1$ not stored)
- 1 sign bit and 11 bits for exponent
- about 16 significant base-10 digits precision with range 10^{-308} to 10^{308}

Arithmetic manipulations

arithmetic with floating-point numbers introduces additional round-off error

- for simplicity: use base-10 numbers, 4-digit mantissa, 1-digit exponent, chopping
- other number bases and rounding would behave in a similar fashion
- focus on: addition, subtraction, multiplication, division
- long sequences of operations can accumulate small round-off errors

Addition

mantissa of no. with smaller exponent is modified so that exponents are the same

Example

$$0.1557 \cdot 10^1 + 0.4381 \cdot 10^{-1}$$

- write

$$0.4381 \cdot 10^{-1} \rightarrow 0.004381 \cdot 10^1$$

- then add

$$\begin{array}{r} 0.1557 \cdot 10^1 \\ 0.004381 \cdot 10^1 \\ \hline 0.160081 \cdot 10^1 \end{array} \Rightarrow \text{chop} \rightarrow 0.1600 \cdot 10^1$$

Adding a large and a small number: add 4000 to 0.0010:

$$0.4000 \cdot 10^4 + 0.0000001 \cdot 10^4 = 0.4000001 \cdot 10^4$$

$$\text{chop} \rightarrow 0.4000 \cdot 10^4$$

Example: MATLAB

sum 0.0001 to itself 10^4 times, which should give 1
in MATLAB running

```
s = 0;  
for i = 1:10000  
s = s + 0.0001;  
end  
sout = s
```

gives

```
sout=0.99999999999991
```

while 0.0001 is a nice number in base-10, it cannot be expressed exactly in base-2

Subtraction

Subtraction

$$0.3641 \cdot 10^2 - 0.2686 \cdot 10^2 = 0.0955 \cdot 10^2$$

$$0.0955 \cdot 10^2 \rightarrow 0.9550 \cdot 10^1 = 9.550$$

zero added to the end is not significant but is appended to fill the empty space

Subtracting two nearly equal numbers

$$0.7642 \cdot 10^3 - 0.7641 \cdot 10^3 = 0.0001 \cdot 10^3$$

$$0.0001 \cdot 10^3 \rightarrow 0.1000 \cdot 10^0 = 0.1000$$

three nonsignificant zeros are appended

round-off induced when subtracting two nearly equal floating-point numbers is called
subtractive cancellation

Multiplication and division

Multiplication: multiply mantissas, add exponents, then normalize and chop

$$(0.1363 \cdot 10^3) \times (0.6423 \cdot 10^{-1}) = 0.08754549 \cdot 10^2$$

$$0.08754549 \cdot 10^2 \rightarrow 0.8754549 \cdot 10^1 \xrightarrow{\text{chop}} 0.8754 \cdot 10^1$$

Division: divide mantissas, subtract exponents, then normalize and chop

Outline

- significant figures
- numerical errors
- round-off errors
- **Taylor series, truncation errors**
- error propagation

Taylor series

on an interval containing x_i and x_{i+1} , we have

$$f(x_{i+1}) = f(x_i) + f'(x_i)h + \frac{f''(x_i)}{2!}h^2 + \cdots + \frac{f^{(n)}(x_i)}{n!}h^n + R_n$$

with $h = x_{i+1} - x_i$ and remainder (in *derivative* or *Lagrange form*)

$$R_n = \frac{f^{(n+1)}(\xi)}{(n+1)!} (x_{i+1} - x_i)^{n+1}$$

for some ξ between x_i and x_{i+1}

- called *Taylor's theorem* or *Taylor series*
- called *Taylor approximation* if R_n is omitted
- provides a polynomial approximation of smooth functions
- predict f at a new point using values and derivatives at a nearby point


Taylor approximations

Zero-order (constant): $f(x_{i+1}) \approx f(x_i)$

First-order (affine approximation): $f(x_{i+1}) \approx f(x_i) + f'(x_i) (x_{i+1} - x_i)$

Second-order (quadratic approximation)

$$f(x_{i+1}) \approx f(x_i) + f'(x_i) (x_{i+1} - x_i) + \frac{f''(x_i)}{2} (x_{i+1} - x_i)^2$$



- approximation improves if x_{i+1} is near x_i
- higher-order terms capture curvature and improve accuracy

Example: polynomial approximation

$$f(x) = -0.1x^4 - 0.15x^3 - 0.5x^2 - 0.25x + 1.2,$$

approximate $f(1) = 0.2$ from $x_i = 0$ with $h = 1$ using $n = 0, 1, 2, 3, 4$

Zero-order: $f(x_{i+1}) \approx f(0) = 1.2$ with $E_t = 0.2 - 1.2 = -1$

First-order: $f'(x) = -0.4x^3 - 0.45x^2 - x - 0.25 \implies f'(0) = -0.25$

$$f(x_{i+1}) \approx 1.2 - 0.25 h \implies f(1) = 0.95$$

$$E_t = 0.2 - 0.95 = -0.75$$

Second-order: $f''(x) = -1.2x^2 - 0.9x - 1 \implies f''(0) = -1$

$$f(x_{i+1}) \approx 1.2 - 0.25 h - \frac{1}{2}h^2 \implies f(1) = 0.45$$

$$E_t = 0.2 - 0.45 = -0.25$$

Fourth-order

$$f^{(3)}(x) = -2.4x - 0.9, \quad f^{(3)}(0) = -0.9, \quad f^{(4)}(x) = -2.4, \quad f^{(4)}(0) = -2.4$$

$$\begin{aligned}f(1) &\approx 1.2 - 0.25(1) - \frac{1}{2}(1)^2 - \frac{0.9}{6}(1)^3 - \frac{2.4}{24}(1)^4 \\&= 1.2 - 0.25 - 0.5 - 0.15 - 0.1 = 0.2\end{aligned}$$

since f is a 4th-degree polynomial, the $n = 4$ Taylor polynomial is exact and $R_4 = 0$

Truncation error order notation

with $h = x_{i+1} - x_i$,

$$R_n = \frac{f^{(n+1)}(\xi)}{(n+1)!} h^{n+1} = O(h^{n+1})$$

- if error is $O(h)$, halving h halves the error
- if $O(h^2)$, halving h quarters the error
- for sufficiently small h , only a few terms are required to get a good estimate

Example

- we use Taylor series to approximate $f(x) = \cos x$ at $x_{i+1} = \pi/3$
- base point $x_i = \pi/4$, step $h = \pi/12$, true value $f(\pi/3) = 0.5$

order n	$f^{(n)}(x)$	$f(\pi/3) \approx$	ε_t
0	$\cos x$	0.707106781	-41.4
1	$-\sin x$	0.521986659	-4.4
2	$-\cos x$	0.497754491	0.449
3	$\sin x$	0.499869147	2.62×10^{-2}
4	$\cos x$	0.500007551	-1.51×10^{-3}
5	$-\sin x$	0.500000304	-6.08×10^{-5}
6	$-\cos x$	0.499999988	2.44×10^{-6}


Remainder and the mean-value insight

- zero-order truncation:

$$f(x_{i+1}) = f(x_i) + R_0, \quad R_0 = f'(x_i)h + \frac{f''(x_i)}{2!}h^2 + \frac{f^{(3)}(x_i)}{3!}h^3 + \dots$$

- derivative mean-value theorem \Rightarrow there exists $\xi \in (x_i, x_{i+1})$ with

$$R_0 = f'(\xi)h$$



Numerical differentiation: forward difference approximation


derivative forward approximation

$$f'(x_i) \approx \frac{f(x_{i+1}) - f(x_i)}{x_{i+1} - x_i}, \quad \text{error } O(x_{i+1} - x_i) = O(h)$$

- follows from Taylor series

$$f(x_{i+1}) = f(x_i) + f'(x_i)(x_{i+1} - x_i) + R_1, \quad R_1 = \frac{f''(\xi)}{2!}(x_{i+1} - x_i)^2$$

- called *finite divided difference*, $f(x_{i+1}) - f(x_i)$ is called *first forward difference*
- $\frac{f(x_{i+1}) - f(x_i)}{x_{i+1} - x_i}$ called *first finite divided difference*




Numerical differentiation: backward difference approximation

derivative backward approximation

$$f'(x_i) \approx \frac{f(x_i) - f(x_{i-1})}{h}, \quad \text{error } O(h)$$

follows from Taylor series expanded backward

$$f(x_{i-1}) = f(x_i) - f'(x_i)h + \frac{f''(x_i)}{2!}h^2 - \dots$$



Numerical differentiation: centered difference approximation

derivative centered difference approximation


$$f'(x_i) = \frac{f(x_{i+1}) - f(x_{i-1})}{2h}, \quad \text{error } O(h^2)$$

follows by subtracting

$$f(x_{i+1}) = f(x_i) + f'(x_i)h + \frac{f''(x_i)}{2!}h^2 - \dots$$

from

$$f(x_{i-1}) = f(x_i) - f'(x_i)h + \frac{f''(x_i)}{2!}h^2 - \dots$$



Example

$$f(x) = -0.1x^4 - 0.15x^3 - 0.5x^2 - 0.25x + 1.25$$

approximate $f'(0.5)$ using $h = 0.5$ and $h = 0.25$ (true value $f'(0.5) = -0.9125$)
for $h = 0.5$, using

$$x_{i-1} = 0, \quad f(x_{i-1}) = 1.2$$

$$x_i = 0.5, \quad f(x_i) = 0.925$$

$$x_{i+1} = 1.0, \quad f(x_{i+1}) = 0.2$$

we get

$$\text{forward: } \frac{0.2 - 0.925}{0.5} = -1.45 \quad (|\varepsilon_t| = 58.9\%)$$

$$\text{backward: } \frac{0.925 - 1.2}{0.5} = -0.55 \quad (|\varepsilon_t| = 39.7\%)$$

$$\text{centered: } \frac{0.2 - 1.2}{1.0} = -1.0 \quad (|\varepsilon_t| = 9.6\%)$$

for $h = 0.25$, using

$$x_{i-1} = 0.25, \quad f(x_{i-1}) = 1.10351563$$

$$x_i = 0.5, \quad f(x_i) = 0.925$$

$$x_{i+1} = 0.75, \quad f(x_{i+1}) = 0.63632813$$

we get

$$\text{forward: } \frac{0.6363 - 0.925}{0.25} = -1.155 \quad (|\varepsilon_t| = 26.5\%)$$

$$\text{backward: } \frac{0.925 - 1.1035}{0.25} = -0.714 \quad (|\varepsilon_t| = 21.7\%)$$

$$\text{centered: } \frac{0.6363 - 1.1035}{0.5} = -0.934 \quad (|\varepsilon_t| = 2.4\%)$$

Numerical differentiation: second-order derivative

Second forward finite divided difference

$$f''(x_i) \approx \frac{f(x_{i+2}) - 2f(x_{i+1}) + f(x_i)}{h^2}, \quad \text{error } O(h)$$

follows by subtracting

$$f(x_{i+2}) = f(x_i) + f'(x_i)(2h) + \frac{f''(x_i)}{2!}(2h)^2 + \dots$$

from 2 times

$$f(x_{i+1}) = f(x_i) + f'(x_i)h + \frac{f''(x_i)}{2!}h^2 + \dots$$

Second backward finite divided difference

$$f''(x_i) \approx \frac{f(x_i) - 2f(x_{i-1}) + f(x_{i-2})}{h^2}, \quad \text{error } O(h)$$

Second centered finite divided difference

$$f''(x_i) \approx \frac{f(x_{i+1}) - 2f(x_i) + f(x_{i-1})}{h^2}, \quad \text{error } O(h^2)$$

Outline

- significant figures
- numerical errors
- round-off errors
- Taylor series, truncation errors
- **error propagation**

Function error propagation

let $\tilde{x} = x + \Delta\tilde{x}$ be an approximation of x ; we seek absolute function error:


$$\Delta f(\tilde{x}) = |f(x) - f(\tilde{x})|$$

- assuming \tilde{x} is close to x and f is continuous and differentiable near x ,

$$f(\tilde{x}) = f(x) + f'(x)\Delta\tilde{x} + \frac{f''(x)}{2} (\Delta\tilde{x})^2 + \dots$$

- dropping second and higher terms:

$$\begin{aligned}\Delta f(\tilde{x}) &\approx |f'(x) (\tilde{x} - x)| \\ &= |f'(x)| |\Delta\tilde{x}|\end{aligned}$$



Example

given $x = 2.5$ with error $\Delta\tilde{x} = 0.01$, estimate the error in $f(x) = x^3$

- we have

$$\Delta f(\tilde{x}) \approx |f'(x)| \Delta\tilde{x} = |3x^2| \cdot 0.01 = 3(2.5)^2(0.01) = 0.1875$$

- since $f(2.5) = 15.625$, we predict

$$f(\tilde{x}) \approx 15.625 \pm 0.1875$$

Stability and condition

Condition (of a problem)

- sensitivity of $f(x)$ to small changes in the input x
- algorithm is *unstable* if it magnifies input/round-off uncertainties
- relative errors:

$$\frac{|f(x) - f(\tilde{x})|}{|f(x)|} \approx \frac{|f'(x) \Delta\tilde{x}|}{|f(x)|}, \quad \frac{|\Delta\tilde{x}|}{|x|} \text{ is the relative error in } x$$

Condition number: ratio of relative errors

$$\kappa(x) = \frac{|x f'(x)|}{|f(x)|}$$

- $\kappa \approx 1$: output relative error \approx input relative error
- $|\kappa| > 1$: relative error is **amplified** (ill-conditioning as $|\kappa|$ grows)
- $|\kappa| < 1$: relative error is **attenuated**

Example

we compute condition number for $f(x) = \tan x$

$$\kappa(x) = \frac{x f'(x)}{f(x)} = \frac{x \sec^2 x}{\tan x} = \frac{x}{\cos^2 x \tan x}$$

Case 1: $x = \frac{\pi}{2} + 0.1\left(\frac{\pi}{2}\right)$

$$\kappa(x) \approx \frac{1.7279 \times 40.86}{-6.314} \approx -11.2$$


Case 2: $x = \frac{\pi}{2} + 0.01\left(\frac{\pi}{2}\right)$

$$\kappa(x) \approx \frac{1.5865 \times 4053}{-63.66} \approx -101$$

near $\pi/2$, $\sec^2 x$ grows and $\tan x$ changes rapidly \Rightarrow severe ill-conditioning

Total numerical error: truncation vs. round-off

- total error = truncation + round-off
- $\downarrow h, \downarrow$ truncation, but \uparrow round-off due to more oper/subtractive cancellation
- goal: choose step size to balance truncation and round-off contributions



Error analysis of numerical differentiation

- centered difference with truncation term:

$$f'(x_i) = \frac{f(x_{i+1}) - f(x_{i-1})}{2h} - \frac{f^{(3)}(\xi)}{6} h^2$$

if the numerator values are exact, the error is due only to truncation

- with rounded values and round-off errors

$$f(x_{i-1}) = \tilde{f}(x_{i-1}) + e_{i-1}, \quad f(x_{i+1}) = \tilde{f}(x_{i+1}) + e_{i+1}$$

- so

$$f'(x_i) = \frac{\tilde{f}(x_{i+1}) - \tilde{f}(x_{i-1})}{2h} + \frac{e_{i+1} - e_{i-1}}{2h} - \frac{f^{(3)}(\xi)}{6} h^2$$

- assume $|e_{i\pm 1}| \leq \varepsilon$ and $|f^{(3)}(x)| \leq M$, then

$$\left| f'(x_i) - \frac{\tilde{f}(x_{i+1}) - \tilde{f}(x_{i-1})}{2h} \right| \leq \frac{\varepsilon}{h} + \frac{M}{6} h^2$$

minimizing the bound w.r.t. h yields the optimal step size $h_{\text{opt}} = (\frac{3\varepsilon}{M})^{1/3}$


Example

- consider $f(x) = -0.1x^4 - 0.15x^3 - 0.5x^2 - 0.25x + 1.2$
- at $x = 0.5$: $f'(0.5) = -0.9125$.
- use centered difference (order $O(h^2)$) with step sizes $h = 10^0, 10^{-1}, \dots, 10^{-10}$

```
function diffex(func, dfunc, x, n)
format long
dftrue = dfunc(x);
h = 1; H(1) = h;
D(1) = (func(x+h) - func(x-h)) / (2*h);
E(1) = abs(dftrue - D(1));
for i = 2:n
    h = h / 10; H(i) = h;
    D(i) = (func(x+h) - func(x-h)) / (2*h);
    E(i) = abs(dftrue - D(i));
end
L = [H' D' E'];
fprintf(' step size    finite difference      true error\n');
fprintf('%14.10f %16.14f %16.13f\n', L);
loglog(H, E), xlabel('Step Size'), ylabel('Error')
title('Plot of Error Versus Step Size')
format short
```

Example

```
>> ff = @(x) -0.1*x.^4 - 0.15*x.^3 - 0.5*x.^2 - 0.25*x + 1.2;  
>> df = @(x) -0.4*x.^3 - 0.45*x.^2 - x - 0.25;  
>> diffex(ff, df, 0.5, 11)
```



- third derivative bound at $x = 0.5$: $M = |f^{(3)}(0.5)| = |-2.4(0.5) - 0.9| = 2.1$
- with machine round-off bound $\varepsilon \approx 0.5 \times 10^{-16}$,

$$h_{\text{opt}} = \left(\frac{3\varepsilon}{M}\right)^{1/3} = \left(\frac{3 \times 0.5 \times 10^{-16}}{2.1}\right)^{1/3} \approx 4.3 \times 10^{-6}$$

References and further readings

- S. C. Chapra and R. P. Canale. *Numerical Methods for Engineers* (8th edition). McGraw Hill, 2021. (Ch.3, Ch.4)
- S. C. Chapra. *Applied Numerical Methods with MATLAB for Engineers and Scientists* (5th edition). McGraw Hill, 2023. (Ch.4)