

12. Algorithms for constrained optimization

- penalty method
- augmented Lagrangian method
- ADMM
- distributed optimization via ADMM

Penalized formulation

$$\begin{array}{ll}\text{minimize} & f(x) \\ \text{subject to} & h_i(x) = 0, \quad i = 1, \dots, p\end{array}$$

Penalized formulation

$$\text{minimize} \quad f(x) + \rho P(h(x))$$

- $h(x) = (h_1(x), \dots, h_p(x))$
- $P : \mathbb{R}^p \rightarrow \mathbb{R}$ is the *penalty function*
- $\rho \in \mathbb{R}$ is the *penalty parameter*
- $\rho P(x)$ penalize constraints violation, *i.e.*, has large values for infeasible points

Penalty function

Penalty function: the penalty function P satisfies the following conditions:

1. P is continuous
2. $P(h(x)) \geq 0$ for all $x \in \mathbb{R}^n$
3. $P(h(x)) = 0$ if and only if x is feasible ($h(x) = 0$)

Quadratic penalty function

$$P(h(x)) = \|h(x)\|^2 = \sum_{i=1}^p (h_i(x))^2$$

Quadratic penalty formulation

$$\text{minimize } f(x) + \rho \|h(x)\|^2$$

- a solution of the above problem might not be feasible
- for large ρ we expect to have small values $(h_i(x))^2$
i.e., an approximate solution to the original problem
- solving the above for an increasing sequence of ρ is called the *penalty method*

Quadratic penalty method

given a starting point $x^{(0)}$, ρ_0 , and a solution tolerance $\epsilon > 0$

repeat for $k = 0, 1, \dots, K$

1. set $x^{(k+1)}$ to be the (approximate) minimizer of

$$\text{minimize } f(x) + \rho_k \|h(x)\|^2$$

using an unconstrained optimization method with initial point $x^{(k)}$

2. update $\rho_{k+1} = 2\rho_k$
-

- terminate if $\|h(x)\|^2$ is small enough
- simple and easy to implement
- but has a major issue:
 - ρ_k rapidly increases with iterations
 - solving penalty problem can be very slow or simply fail

Inequality constraints

$$\begin{array}{ll}\text{minimize} & f(x) \\ \text{subject to} & g_i(x) \leq 0, \quad i = 1, \dots, m \\ & h_j(x) = 0, \quad j = 1, \dots, p\end{array}$$

can be handled using the penalized formulation

$$\text{minimize} \quad f(x) + \rho \|h(x)\|^2 + \rho \|g^+(x)\|^2$$

- $g^+(x) = (g_1^+(x), \dots, g_m^+(x))$ and

$$g_i^+(x) = \max\{0, g_i(x)\} = \begin{cases} 0 & \text{if } g_i(x) \leq 0 \\ g_i(x) & \text{if } g_i(x) > 0 \end{cases}$$

- there are other choices of penalty functions
- we just consider the simple quadratic penalization function

Outline

- penalty method
- **augmented Lagrangian method**
- ADMM
- distributed optimization via ADMM

Original problem

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & h(x) = 0 \end{array} \quad (12.1)$$

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$ and $h : \mathbb{R}^p \rightarrow \mathbb{R}$
- Lagrangian: $L(x, \lambda) = f(x) + \lambda^T h(x)$ where $\lambda \in \mathbb{R}^p$
- problem is equivalent to (for any λ)

$$\begin{array}{ll} \text{minimize} & L(x, \lambda) = f(x) + \lambda^T h(x) \\ \text{subject to} & h(x) = 0 \end{array}$$

- if x^\star is a solution and a regular point, then

$$\nabla_x L(x^\star, \lambda^\star) = 0 \quad \text{for some } \lambda^\star$$

Augmented Lagrangian formulation

the **augmented Lagrangian** (AL) is

$$\begin{aligned} L_{\rho}(x, \lambda) &= L(x, \lambda) + \rho \|h(x)\|^2 \\ &= f(x) + \lambda^T h(x) + \rho \|h(x)\|^2 \end{aligned}$$

- this is the Lagrangian $L(x, \lambda)$ augmented with a quadratic penalty
- ρ is a positive penalty parameter
- augmented Lagrangian is the Lagrangian of the equivalent problem

$$\begin{array}{ll} \text{minimize} & f(x) + \rho \|h(x)\|^2 \\ \text{subject to} & h(x) = 0 \end{array}$$

- solution of the original problem is also a solution of the AL formulation
- AL method minimizes $L_{\rho}(x, \lambda)$ for a sequence of values of λ and ρ

Lagrange multiplier update

- minimizer \tilde{x} of augmented Lagrangian $L_\rho(x, \lambda)$ satisfies

$$\nabla f(\tilde{x}) + Dh(\tilde{x})^T(2\rho h(\tilde{x}) + \lambda) = 0$$

- if we define $\tilde{\lambda} = \lambda + 2\rho h(\tilde{x})$ this can be written as

$$\nabla f(\tilde{x}) + Dh(\tilde{x})^T \tilde{\lambda} = 0$$

- this is the first equation in the optimality conditions

$$\nabla f(\hat{x}) + Dh(\hat{x})^T \hat{\lambda} = 0, \quad h(\hat{x}) = 0$$

- shows that if $h(\tilde{x}) = 0$, then \tilde{x} is optimal
- if $h(\tilde{x})$ is not small, suggests $\tilde{\lambda}$ is a good update for λ

Augmented Lagrangian algorithm

given $x^{(0)}$, $\lambda^{(0)}$, $\rho^{(0)}$, and a solution tolerance $\epsilon > 0$

repeat for $k = 0, 1, \dots$

1. set $x^{(k+1)}$ to be an (approximate) solution to

$$\text{minimize } f(x) + (\lambda^{(k)})^T h(x) + \rho^{(k)} \|h(x)\|^2$$

using any unconstrained optimization method with initial point $x^{(k)}$

2. update $\lambda^{(k)}$:

$$\lambda^{(k+1)} = \lambda^{(k)} + 2\rho^{(k)} h(x^{(k+1)})$$

3. set $\rho^{(k)}$ as constant or

$$\begin{cases} \rho^{(k)} & \text{if } \|h(x^{(k+1)})\| < 0.25\|h(x^{(k)})\| \\ 2\rho^{(k)} & \text{otherwise} \end{cases}$$

-
- ρ is increased only when needed, more slowly than in penalty method
 - continues until $h(x^{(k)})$ and/or $\nabla L(x^{(k)}, \lambda^{(k)})$ are sufficiently small

Example

consider applying the augmented Lagrangian method to the problem:

$$\begin{array}{ll}\text{minimize} & e^{3x_1} + e^{-4x_2} \\ \text{subject to} & x_1^2 + x_2^2 = 1\end{array}$$

with $x^{(0)} = (-1, 1)$ and $\lambda^{(0)} = -1$, we set a constant penalty parameter $\rho_k = 10$

the augmented Lagrangian function is

$$L_\rho(x, \lambda) = e^{3x_1} + e^{-4x_2} + \lambda (x_1^2 + x_2^2 - 1) + (\rho/2) (x_1^2 + x_2^2 - 1)^2$$

for the inner minimization problems, we employ Newton's method:

$$\hat{x} \leftarrow \hat{x} + \nabla^2 L_\rho(\hat{x}, \lambda^{(k)})^{-1} \nabla L_\rho(\hat{x}, \lambda^{(k)})$$

the gradient and Hessian are:

$$\nabla L_{\rho}(x, \lambda) = \begin{bmatrix} 3e^{3x_1} + 2\lambda x_1 + 2\rho x_1(x_1^2 + x_2^2 - 1) \\ -4e^{-4x_2} + 2\lambda x_2 + 2\rho x_2(x_1^2 + x_2^2 - 1) \end{bmatrix}$$

and

$$\nabla^2 L_{\rho}(x, \lambda) = \begin{bmatrix} 9e^{3x_1} + 2\lambda + 2\rho(x_1^2 + x_2^2 - 1) + 4\rho x_1^2 & 4\rho x_1 x_2 \\ 4\rho x_1 x_2 & 16e^{-4x_2} + 2\lambda + 2\rho(x_1^2 + x_2^2 - 1) + 4\rho x_2^2 \end{bmatrix}$$

iteration starts from $\hat{x} = x^{(k)}$ and continues until $\|\nabla L_{\rho}(\hat{x}, \lambda^{(k)})\| < 10^{-4}$

the value $x^{(k+1)}$ is then set to \hat{x} and the Lagrange multiplier is subsequently updated:

$$\lambda^{(k+1)} = \lambda^{(k)} + \rho((x_1^{(k+1)})^2 + (x_2^{(k+1)})^2 - 1)$$

MATLAB code implementation

```
rho=10;
x=[-1;1];
lam=-1;
%% AL gradient and Hessian
g=@(x,lam)[3*exp(3*x(1))+2*lam*x(1)+2*rho*x(1)*(x(1)^2+x(2)^2-1);
-4*exp(-4*x(2))+2*lam*x(2)+2*rho*x(2)*(x(1)^2+x(2)^2-1)];
hess=@(x,lam)[9*exp(3*x(1))+2*lam+2*rho*(x(1)^2+x(2)^2-1)+4*rho*x(1)^2 4*rho*x(1)*x(2);
4*rho*x(1)*x(2) 16*exp(-4*x(2))+2*lam+2*rho*(x(1)^2+x(2)^2-1)+4*rho*x(2)^2];
%% AL method
for i=1:100
% Newton inner minimization
while (norm(g(x,lam)) >= 1e-4)
d = -hess(x,lam)\g(x,lam);
x = x+d;
end
% Lagrange update
lam=lam+rho*(x(1)^2+x(2)^2-1);
end
```

running for 100 iterations, we get $x^* = (-0.7483, 0.6633)$ and $\lambda^* = 0.2123$

Outline

- penalty method
- augmented Lagrangian method
- **ADMM**
- distributed optimization via ADMM

ADMM problem form

the alternating direction method of multiplier (ADMM) solves problem of form:

$$\begin{array}{ll}\text{minimize} & f(x) + g(z) \\ \text{subject to} & Ax + Bz = c\end{array}$$

- variables $x \in \mathbb{R}^n$ and $z \in \mathbb{R}^m$
- $A \in \mathbb{R}^{p \times n}$, $B \in \mathbb{R}^{p \times m}$, and $c \in \mathbb{R}^p$
- augmented Lagrangian

$$L_\rho(x, z, \lambda) = f(x) + g(z) + \lambda^T(Ax + Bz - c) + (\rho/2)\|Ax + Bz - c\|^2$$

ADMM update

$$x^{(k+1)} = \underset{x}{\operatorname{argmin}} L_{\rho}(x, z^{(k)}, \lambda^{(k)})$$

$$z^{(k+1)} = \underset{z}{\operatorname{argmin}} L_{\rho}(x^{(k+1)}, z, \lambda^{(k)})$$

$$\lambda^{(k+1)} = \lambda^{(k)} + \rho(Ax^{(k+1)} + Bz^{(k+1)} - c)$$

- $\rho > 0$ is the ADMM penalty parameter
- x and z are updated in an alternating or sequential fashion
- this is different from AL method where x and z are minimized jointly

$$(x, z) = \underset{x, z}{\operatorname{argmin}} L_{\rho}(x, z, \lambda^{(k)})$$

- separating the minimization over x and z allows to decompose large problems into smaller ones when f or g are separable

ADMM scaled form

define the residual $r = Ax + Bz - c$ and $u = (1/\rho)\lambda$, then

$$\begin{aligned}\lambda^T r + (\rho/2)\|r\|^2 &= (\rho/2)\|r + (1/\rho)\lambda\|^2 - (1/2\rho)\|\lambda\|^2 \\ &= (\rho/2)\|r + u\|^2 - (\rho/2)\|u\|^2\end{aligned}$$

ADMM scaled form

$$\begin{aligned}x^{(k+1)} &= \underset{x}{\operatorname{argmin}} \left(f(x) + (\rho/2)\|Ax + Bz^{(k)} - c + u^{(k)}\|^2 \right) \\ z^{(k+1)} &= \underset{z}{\operatorname{argmin}} \left(g(z) + (\rho/2)\|Ax^{(k+1)} + Bz - c + u^{(k)}\|^2 \right) \\ u^{(k+1)} &= u^{(k)} + Ax^{(k+1)} + Bz^{(k+1)} - c\end{aligned}$$

Example: quadratic programs

$$\begin{array}{ll}\text{minimize} & (1/2)x^T Q x + r^T x \\ \text{subject to} & Cx = d \\ & x \geq 0\end{array}$$

- Q is positive semidefinite (reduces to an LP when $Q = 0$)
- we can express this problem in the ADMM form:

$$\begin{array}{ll}\text{minimize} & f(x) + g(z) \\ \text{subject to} & x - z = 0\end{array}$$

where

$$f(x) = (1/2)x^T Q x + r^T x, \quad \text{dom } f = \{x \mid Cx = d\}$$

is the original objective with restricted domain

- g is the indicator function of the nonnegative orthant \mathbb{R}_+^n

the scaled form of ADMM consists of the iterations

$$\begin{aligned}x^{(k+1)} &= \underset{x}{\operatorname{argmin}} \left(f(x) + (\rho/2) \|x - z^{(k)} + u^{(k)}\|^2 \right) \\z^{(k+1)} &= (x^{(k+1)} + u^{(k)})_+ \\u^{(k+1)} &= u^{(k)} + x^{(k+1)} - z^{(k+1)}\end{aligned}$$

the x -update is a constrained least squares problem with optimality conditions

$$\begin{bmatrix} Q + \rho I & C^T \\ C & 0 \end{bmatrix} \begin{bmatrix} x^{(k+1)} \\ v \end{bmatrix} + \begin{bmatrix} r - \rho(z^{(k)} - u^{(k)}) \\ -d \end{bmatrix} = 0$$

Norm-one regularized least squares

the **lasso** problem is the ℓ_1 regularized least squares

$$\text{minimize} \quad (1/2)\|Ax - b\|^2 + \eta\|x\|_1$$

- $\eta > 0$ is a scalar regularization parameter
- in ADMM form, the lasso problem can be written as

$$\begin{array}{ll}\text{minimize} & f(x) + g(z) \\ \text{subject to} & x - z = 0\end{array}$$

where $f(x) = (1/2)\|Ax - b\|^2$ and $g(z) = \eta\|z\|_1$

the ADMM iteration is

$$x^{(k+1)} = (A^T A + \rho I)^{-1} (A^T b + \rho(z^{(k)} - u^{(k)}))$$

$$z^{(k+1)} = S_{\eta/\rho}(x^{(k+1)} + u^{(k)})$$

$$u^{(k+1)} = u^{(k)} + x^{(k+1)} - z^{(k+1)}$$

where S is the soft thresholding operator defined element-wise as

$$\begin{aligned} S_{\kappa}(a) &= \begin{cases} a - \kappa & a > \kappa \\ 0 & |a| \leq \kappa \\ a + \kappa & a < -\kappa \end{cases} \\ &= (a - \kappa)_+ - (-a - \kappa)_+ \end{aligned}$$

Outline

- penalty method
- augmented Lagrangian method
- ADMM
- **distributed optimization via ADMM**

Consensus problem

$$\text{minimize } f(x) = \sum_{i=1}^N f_i(x)$$

- variable $x \in \mathbb{R}^n$
- $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ represents the i th component of the objective function
- f_i is available only on machine processor i
- goal is to solve this problem with f_i handled by processor i only

Example

many classification or regression problems can be formulated as:

$$\text{minimize} \quad \sum_{j=1}^m l(x; \xi_j)$$

- $l(x; \xi_j)$ represent the loss function for data ξ_j
- for large m , storing the data on a single machine may not be feasible
- the problem can be solved by distributing the data across multiple machines,

$$f_i(x) = \sum_{j \in \mathcal{J}_i} l(x; \xi_j),$$

where \mathcal{J}_i is the set of training data indices at machine i

Equivalent formulation

$$\begin{array}{ll}\text{minimize} & \sum_{i=1}^N f_i(x_i) \\ \text{subject to} & x_i - z = 0, \quad i = 1, \dots, N\end{array}$$

- $x_i \in \mathbb{R}^n$ handled by processing unit i
- z is a global variable z handled by central processing unit called *central server*
- the constraints ensure that all local variables are equal
- objective is now separable in the variables x_i
- the augmented Lagrangian is

$$L_\rho(x_1, \dots, x_N, z, \lambda) = \sum_{i=1}^N \left(f_i(x_i) + (\lambda_i)^T (x_i - z) + \frac{\rho}{2} \|x_i - z\|^2 \right)$$

ADMM updates

$$x_i^{(k+1)} = \underset{x_i}{\operatorname{argmin}} (f_i(x_i) + \lambda_i^{(k)T}(x_i - z^{(k)}) + \frac{\rho}{2} \|x_i - z^{(k)}\|^2)$$

$$z^{(k+1)} = \frac{1}{N} \sum_{i=1}^N (x_i^{(k+1)} + \frac{1}{\rho} \lambda_i^{(k)})$$

$$\lambda_i^{(k+1)} = \lambda_i^{(k)} + \rho(x_i^{(k+1)} - z^{(k+1)})$$

- the first and last steps are updated independently by each machine i
- central server updates z after it receives all x_i and then send it back to machines

Equivalent simpler update

using overline to denote the average of a vector, we can express the z -update as:

$$z^{(k+1)} = \bar{x}^{(k+1)} + \frac{1}{\rho} \bar{\lambda}^{(k)}$$

by taking the average of the λ -update, we get:

$$\bar{\lambda}^{(k+1)} = \bar{\lambda}^{(k)} + \rho(\bar{x}^{(k+1)} - z^{(k+1)})$$

substituting the first equation into the subsequent one, we obtain $\bar{\lambda}^{(k+1)} = 0$ for all k

hence $z^{(k)} = \bar{x}^{(k)}$ and ADMM can be reformulated as:

$$\begin{aligned} x_i^{(k+1)} &= \underset{x_i}{\operatorname{argmin}} (f_i(x_i) + \lambda_i^{(k)T} (x_i - \bar{x}^{(k)}) + \frac{\rho}{2} \|x_i - \bar{x}^{(k)}\|^2) \\ \lambda_i^{(k+1)} &= \lambda_i^{(k)} + \rho(x_i^{(k+1)} - \bar{x}^{(k+1)}) \end{aligned}$$

Regularized consensus problem

$$\begin{aligned} & \text{minimize} && \sum_{i=1}^N f_i(x_i) + g(z) \\ & \text{subject to} && x_i - z = 0, \quad i = 1, \dots, N, \end{aligned}$$

- objective term g is a constraint or regularization (e.g., $g(z) = \|z\|_1$)
- for this case, the ADMM method is:

$$x_i^{(k+1)} = \underset{x_i}{\operatorname{argmin}} (f_i(x_i) + \lambda_i^{(k)T}(x_i - z^{(k)}) + \frac{\rho}{2} \|x_i - z^{(k)}\|^2)$$

$$z^{(k+1)} = \underset{z}{\operatorname{argmin}} (g(z) + \sum_{i=1}^N (-\lambda_i^{(k)T} z + \frac{\rho}{2} \|x_i^{(k+1)} - z\|^2))$$

$$\lambda_i^{(k+1)} = \lambda_i^{(k)} + \rho(x_i^{(k+1)} - z^{(k+1)})$$

- collecting linear and quadratic terms, the z -update can be expressed as:

$$z^{(k+1)} = \underset{z}{\operatorname{argmin}} \left(g(z) + \frac{N\rho}{2} \|z - \bar{x}^{(k+1)} - \frac{1}{\rho} \bar{\lambda}^{(k)}\|^2 \right)$$

- when g is nonzero, we don't typically get that $\bar{\lambda}^{(k)} = 0$
- hence λ_i terms cannot be eliminated as in the non-regularized case
- using the above update form for z , ADMM is:

$$x_i^{(k+1)} = \underset{x_i}{\operatorname{argmin}} \left(f_i(x_i) + \lambda_i^{(k)T} (x_i - z^{(k)}) + \frac{\rho}{2} \|x_i - z^{(k)}\|^2 \right)$$

$$z^{(k+1)} = \underset{z}{\operatorname{argmin}} \left(g(z) + \frac{N\rho}{2} \|z - \bar{x}^{(k+1)} - \frac{1}{\rho} \bar{\lambda}^{(k)}\|^2 \right)$$

$$\lambda_i^{(k+1)} = \lambda_i^{(k)} + \rho (x_i^{(k+1)} - z^{(k+1)})$$

Examples

- for $g(z) = \eta \|z\|_1$, the z -update translates into a soft threshold operation:

$$z^{(k+1)} = S_{\eta/N\rho}(\bar{x}^{(k+1)} - \frac{1}{\rho}\bar{\lambda}^{(k)})$$

- considering g as the indicator function of \mathbb{R}_+^n , then

$$z^{(k+1)} = (\bar{x}^{(k+1)} - \frac{1}{\rho}\bar{\lambda}^{(k)})_+$$

for this problem, the scaled variant of ADMM, exhibited below, is often more streamlined and manageable compared to its unscaled counterpart:

$$x_i^{(k+1)} = \underset{x_i}{\operatorname{argmin}} (f_i(x_i) + \frac{\rho}{2} \|x_i - z^{(k)} + u_i^{(k)}\|^2)$$

$$z^{(k+1)} = \underset{z}{\operatorname{argmin}} (g(z) + \frac{N\rho}{2} \|z - \bar{x}^{(k+1)} - \bar{u}^{(k)}\|^2)$$

$$u_i^{(k+1)} = u_i^{(k)} + x_i^{(k+1)} - z^{(k+1)}$$

References and further readings

- I. Griva and S. G. Nash and A. Sofer. *Linear and Nonlinear Optimization*. SIAM, 2009.
- E. K.P. Chong, Wu-S. Lu, and S. H. Zak. *An Introduction to Optimization: With Applications to Machine Learning*. John Wiley & Sons, 2023. (chapter 14)
- S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein. *Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers*. Foundations and Trends in Machine learning, 2011.